

Lecture 10: Graph Search

Harvard SEAS - Fall 2022

2022-10-06

1 Announcements

Recommended Reading:

- Roughgarden II Sec 8.1–8.2
- CLRS 22.2

2 Shortest Walks

Motivated by a (simplified version) of the Google Maps problem, we wish to design an algorithm for the following computational problem:

<p>Input : A digraph $G = (V, E)$ and two vertices $s, t \in V$</p> <p>Output : A <i>shortest walk</i> from s to t in G, if any walk from s to t exists</p>
--

Computational Problem ShortestWalk

Definition 2.1. Let $G = (V, E)$ be a directed graph, and $s, t \in V$.

- A *walk* w from s to t in G is a sequence v_0, v_1, \dots, v_ℓ of vertices such that $v_0 = s$, $v_\ell = t$, and $(v_{i-1}, v_i) \in E$ for $i = 1, \dots, \ell$.
- The *length* of a walk w is $\text{length}(w) =$ the number of edges in w (the number ℓ above).
- The *distance* from s to t in G is

$$\text{dist}_G(s, t) = \begin{cases} \min\{\text{length}(w) : w \text{ is a walk from } s \text{ to } t\} & \text{if a walk exists} \\ \infty & \text{otherwise} \end{cases}$$

- A *shortest walk* from s to t in G is a walk w from s to t with $\text{length}(w) = \text{dist}_G(s, t)$

Q: An algorithm immediate from the definition?

A: Enumerate over all walks from s in order of length, and terminate after finding the first that ends at t .

But when can we stop this algorithm to conclude that there is no walk? The following lemma allows us to stop at walks of length $n - 1$.

Lemma 2.2. *If w is a shortest walk from s to t , then all of the vertices that occur on w are distinct. That is, every shortest walk is a path — a walk in which all vertices are distinct.*

Proof.

□

Q: With this lemma, what is the runtime of exhaustive search?

A:

3 Breadth-First Search

We can get a faster algorithm using *breadth-first search (BFS)*. For simplicity we'll start by presenting the algorithm for the following simpler computational problem:

Input : A directed graph $G = (V, E)$ and two vertices $s, t \in V$ Output : The distance from s to t in G

Computational Problem DistanceInGraph

<pre>1 BFSv0(G, s, t) Input : A directed graph $G = (V, E)$ and two vertices $s, t \in V$ Output : The distance from s to t in G 2 $S = \{s\};$ 3 <i>/* loop invariant: */</i> 4 foreach $d = 0, \dots, n - 1$ do 5 if $t \in S$ then return $d;$ 6 $S = S \cup \{v \in V : \exists u \in S \text{ s.t. } (u, v) \in E\}$ 7 return ∞</pre>
--

Example:

Q: What is happening at every iteration of the loop?

We have a set of S which is the set of vertices that have been visited previously. At each iteration, we construct a new S' that is the union of S and the set of vertices that can be visited from all the vertices in S by one additional edge. This allows us to include the new vertices that can be visited now that we update the distance d .

Q: How do we we perform the update of Line 6?

We'll iterate over all edges of G . We assume we are given the graph as an *adjacency list*: for each vertex v , we keep a neighbor array $\text{Nbr}[v] = \{u : (v, u) \in E\}$ holding the neighbors of v . We are also given the length of each such array $\text{Nbr}[v]$ —we could compute these lengths ourselves, but they're so often useful that we'll save time by assuming the representation of the graph comes with them.

¹ So we iterate over all edges of G by iterating over all those lists.

Q: How do we prove correctness?

Q: What is the runtime of the algorithm, in terms of the number of vertices n and the number of edges m ?

4 Improving BFS

Observations:

- S only grows due to edges that cross the *frontier* from S to $V - S$.
- Every edge in E crosses the frontier in at most one loop iteration.

¹Other ways of representing a graph are sometimes useful, and discussed in classes like CS 124. In CS 120, we'll always represent graphs by adjacency lists.

```

1 BFS( $G, s, t$ )
   Input    : A directed graph  $G = (V, E)$  and two vertices  $s, t \in V$ 
   Output   : The distance from  $s$  to  $t$  in  $G$ 
2  $S = \{s\}$ ;
3  $F = \{s\}$  ;                               /* the frontier vertices */
4  $d = 0$ ;
5 /* loop invariant:      */
6 while  $F \neq \emptyset$  do
7   | if  $t \in F$  then return  $d$ ;
8   |  $F = \{v \in V - S : \exists u \in F \text{ s.t. } (u, v) \in E\}$ ;
9   |  $S = S \cup F$ ;
10  |  $d = d + 1$ ;
11 return  $\infty$ 

```

Theorem 4.1. *BFS(G) correctly solves DistanceInGraph and can be implemented in time $O(n+m)$, where n is the number of vertices in G and m is the number of edges.*

Proof. 1. Correctness:

2. Runtime:

□

Above we used the following definition:

Definition 4.2. For a digraph $G = (V, E)$ and a vertex v , we define the *out-degree* of v to be

$$d_{out}(v) = |\{w : (v, w) \in E\}|$$

and the *in-degree* of v to be

$$d_{in}(v) = |\{u : (u, v) \in E\}|.$$

For an undirected graph, we have $d_{out}(v) = d_{in}(v)$, so we just call this the *degree* of v , denoted $d(v)$.

Q: How would we calculate the out-degree of v from the adjacency-list representation of a graph?

5 More Graph Search

Q: How to actually find a shortest *path*, not just the distance?

Note that, by Lemma 2.2, shortest walks are paths, so we can use the terms “shortest paths” and “shortest walks” interchangeably.

Observation: BFS actually solves the following computational problem:

Input : A digraph $G = (V, E)$ and a vertex $s \in V$
Output : For every vertex v , $\text{dist}_G(s, v)$ and, if $\text{dist}_G(s, v) < \infty$, a path p_v from s to v of length $\text{dist}_G(s, v)$ (implicitly represented through a predecessor array as above)

Computational Problem SingleSourceShortestPaths

We have proven:

Theorem 5.1. *There is an algorithm that solves SingleSourceShortestPaths in time $O(n + m)$ on digraphs with n vertices and m edges in adjacency list representation.*

The algorithm we have seen (BFS) only works on unweighted graphs; algorithms for weighted graphs are covered in CS124.

6 Other Forms of Graph Search

Another very useful form of graph search that you may have seen is *depth-first search* (DFS). We won't cover it in CS120, but DFS and some of its applications are covered in CS124.

We do, however, briefly mention a randomized form of graph search, namely *random walks*, and use it to solve the *decision* problem of STConnectivity on undirected graphs.

Input : A graph $G = (V, E)$ and vertices $s, t \in V$ Output : YES if there is a walk from s to t in G , and NO otherwise

Computational Problem UndirectedSTconnectivity

<pre>1 RandomWalk(G, s, ℓ) Input : A digraph $G = (V, E)$, a vertices $s, t \in V$, and a walk-length ℓ Output : YES or NO 2 $v = s$; 3 foreach $i = 1, \dots, \ell$ do 4 if $v = t$ then return YES; 5 $j = \text{random}(d_{out}(v))$; 6 $v = j$'th out-neighbor of v; 7 return ∞</pre>
--

Q: What is the advantage of this algorithm over BFS?

It can be shown that if G is an *undirected* graph with n vertices and m edges, then for an appropriate choice of $\ell = O(mn)$, with high probability $\text{RandomWalk}(G, s, \ell)$ will visit all vertices reachable from s . Thus, we obtain a *Monte Carlo* algorithm for UndirectedSTConnectivity.

Theorem 6.1. *UndirectedSTConnectivity can be solved by a Monte Carlo randomized algorithm with arbitrarily small error probability in time $O(mn)$ using only $O(1)$ words of memory in addition to the input.*